

FakeNewsNet: A Data Repository with News Content, Social Context and Dynamic Information for Studying Fake News on Social Media

Kai Shu¹, Deepak Mahudeswaran¹, Suhang Wang², Dongwon Lee² and Huan Liu¹

¹Computer Science and Engineering, Arizona State University, Tempe, 85281, USA

²College of Information Sciences and Technology, Penn State University, University Park, PA, 16802, USA
{kai.shu, dmahudes, huan.liu}@asu.edu, {szw494, dlee}@ist.psu.edu

Abstract

Social media has become a popular means for people to consume news. Meanwhile, it also enables the wide dissemination of *fake news*, i.e., news with intentionally false information, which brings significant negative effects to the society. Thus, fake news detection is attracting increasing attention. However, fake news detection is a non-trivial task, which requires multi-source information such as *news content*, *social context*, and *dynamic information*. First, fake news is written to fool people, which makes it difficult to detect fake news simply based on news contents. In addition to news contents, we need to explore social contexts such as user engagements and social behaviors. For example, a credible user's comment that "this is a fake news" is a strong signal for detecting fake news. Second, dynamic information such as how fake news and true news propagate and how users' opinions toward news pieces are very important for extracting useful patterns for (early) fake news detection and intervention. Thus, comprehensive datasets which contain *news content*, *social context*, and *dynamic information* could facilitate fake news propagation, detection, and mitigation; while to the best of our knowledge, existing datasets only contains one or two aspects. Therefore, in this paper, to facilitate fake news related researches, we provide a fake news data repository *FakeNewsNet*, which contains two comprehensive datasets that includes *news content*, *social context*, and *dynamic information*. We present a comprehensive description of datasets collection, demonstrate an exploratory analysis of this dataset from different perspectives, and discuss the benefits of FakeNewsNet for potential applications on fake news study on social media.

1 Introduction

Social media has become a primary source of news consumption nowadays. Social media is free of cost, easy to access and helps one to express opinions publicly and hence acts as an excellent way for individuals to consume information. For example, the time individuals spend on social media is continually increasing¹. As another example, studies from Pew Research Center shows that around 67% of Amer-

icans get some of their news on social media² and this have shown constant increase since 2016. Since there is no regulatory authority on social media, the quality of news pieces spread in social media is often lower than traditional news sources. In other words, social media also enable the wide spread of fake news. Fake news (Shu et al. 2017), as a specific type of disinformation, means the false information that is spread deliberately to deceive people. Fake news affects the individuals as well as society as a whole. First, fake news can disturb the authenticity balance of the news ecosystem. Second, fake news persuades consumers to accept false or biased stories. For example, some individuals and organizations spread fake news in social media for financial and political gains (Shu et al. 2017). It is also reported that fake news has influence on the 2016 US presidential elections³. Finally, fake news changes the way people interpret and respond to real news. Thus, fake news detection is a critical issue that need to be addressed.

Detecting fake news on social media present unique challenges. First, fake news pieces are intentionally written to mislead consumers, which makes it not satisfactory to spot fake news from news content itself. Thus, we need to explore information in addition to news content, such as social engagements and social behaviors of users on social media. For example, a credible user's comment that "This is bull shit" is a strong signal that the user doesn't believe the news and thus the news may be fake. Second, the research community lacks datasets which contain dynamic information to understand how fake news propagates, how users react to fake news, and how we can extract useful temporal patterns for (early) fake news detection and intervention. Thus, it is necessary to have comprehensive datasets that have news content, social context and dynamic information to facilitate fake news research. However, to the best of our knowledge, existing datasets only covers one or two aspects.

Therefore, in this paper, we construct and public a multi-dimension data repository *FakeNewsNet*, which contains two datasets with news content, social context, and dynamic information. The constructed FakeNewsNet repository has

the potential to boost the study of various open research problems related to fake news study. First, the rich set of features in the datasets provides an opportunity to experiment with different approaches for fake news detection and understand the diffusion of fake news in social network and intervene in it. Second, the dynamic information enables to study early fake news detection by generating synthetic user engagements from historical temporal user engagement patterns in the dataset (Qian et al. 2018a). Third, we can investigate the fake news diffusion process by identifying provenances, persuaders, and developing better fake news intervention strategies (Shu, Bernard, and Liu 2018). Our data repository can serve as a starting point for many exploratory studies for fake news, and provide a better, shared insight into disinformation tactics. We aim to continuously update this data repository, expand it with new sources and features, as well as maintain completeness.

The rest of the paper is organized as follows. We describe the background of fake news problem and the existing related datasets in Section 2. In Section 3, we introduce the details of data construction. We further perform exploration analysis on FakeNewsNet and provide some insights from different perspectives to study fake news in Section 4. In Section 5, we discuss the potential applications of the data repository, and we conclude with future work in Section 6.

2 Background and Related Work

Fake news detection in social media aims to extract useful features and build effective models from existing social media datasets for detecting fake news in the future. Thus, a comprehensive and large-scale dataset with multi-dimension information in online fake news ecosystem is important. The multi-dimension information not only provides more signals for detecting fake news, but can also be used for researches such as understanding fake news propagation and fake news intervention. Though there exist several datasets for fake news detection, the majority of them only contains linguistic features. Few of them contains both linguistic and social context features. To facilitate research on fake news, we provide a data repository which include not only news contents and social contents, but also dynamic information. For a better comparison of the differences, we list existing popular fake news detection datasets below and compare them with *FakeNewsNet* in Table 2.

BuzzFeedNews⁴: This dataset comprises a complete sample of news published in Facebook from 9 news agencies over a week close to the 2016 U.S. election from September 19 to 23 and September 26 and 27. Every post and the linked article were fact-checked claim-by-claim by 5 BuzzFeed journalists. It contains 1,627 articles 826 mainstream, 356 left-wing, and 545 right-wing articles.

LIAR⁵: This dataset (Wang 2017) is collected from fact checking website PolitiFact. It has 12.8 K human labeled short statements collected from PolitiFact and the statements are labeled into six categories ranging from completely false

to completely true as pants on fire, false, barely-true, half-true, mostly true, and true.

BS Detector⁶: This dataset is collected from a browser extension called BS detector developed for checking news veracity. It searches all links on a given web page for references to unreliable sources by checking against a manually compiled list of domains. The labels are the outputs of BS detector, rather than human annotators.

CRED BANK⁷: This is a large-scale crowd-sourced dataset (Mittra and Gilbert 2015) of around 60 million tweets that cover 96 days starting from October 2015. All the tweets are related to over 1,000 news events, with each event assessed for credibilities by 30 annotators from Amazon Mechanical Turk.

BuzzFace⁸: This dataset (Santia and Williams 2018) collected using by extending the BuzzFeed dataset with comment related to news articles in Facebook. The dataset contains 2263 news articles and 1.6 million comments discussing news content.

FacebookHoax⁹: This dataset (Tacchini et al. 2017) comprises information related to posts from the Facebook pages related to scientific news (non- hoax) and conspiracy pages (hoax) collected using Facebook Graph API. The dataset contains 15,500 posts from 32 pages (14 conspiracy and 18 scientific) with more than 2,300,000 likes.

From Table 2, we observe that no existing public dataset can provide all possible features of interests. Existing datasets have some limitation that we try to address in our data repository. For example, BuzzFeedNews only contains headlines and text for each news piece and covers news articles from very few news agencies. LIAR dataset contains mostly short statements instead of entire news articles with the meta attributes. BS Detector data is collected and annotated by using a developed news veracity checking tool, rather than using human expert annotators. CRED BANK dataset was originally collected for evaluating tweet credibilities and the tweets in the dataset are not related to the fake news articles and hence cannot be effectively used for fake news detection. BuzzFace dataset has basic news contents and social context information but it does not capture the changing dynamic information. The FacebookHoax dataset consists very few instances about the conspiracy theories and scientific news.

To address the disadvantages of existing fake news detection dataset, the proposed FakeNewsNet repository collects multi-dimension information from news content, social context, and dynamic information from different types of news domains such as political and entertainment sources.

3 Dataset Construction

In this section, we introduce the dataset construction process for FakeNewsNet repository. We demonstrate how we can collect news contents with reliable ground truth labels, how we obtain additional social context information, and

⁴<https://github.com/BuzzFeedNews/2016-10-facebook-fact-check/tree/master/data>

⁵<https://www.cs.ucsb.edu/~william/software.html>

⁶<https://github.com/bs-detector/bs-detector>

⁷<http://compsocial.github.io/CREDBANK-data/>

⁸<https://github.com/gsantia/BuzzFace>

⁹<https://github.com/gabll/some-like-it-hoax>

Table 1: Comparison with existing fake news detection datasets

Dataset \ Features	News Content		Social Context				Dynamic Information
	Linguistic	Visual	User	Post	Second order	Network	
BuzzFeedNews	✓	✗	✗	✗	✗	✗	✗
LIAR	✓	✗	✗	✗	✗	✗	✗
BS Detector	✓	✗	✗	✗	✗	✗	✗
CREDBANK	✓	✗	✓	✓	✗	✓	✗
BuzzFace	✓	✗	✗	✓	✓	✗	✗
FacebookHoax	✓	✗	✓	✓	✓	✗	✗
FakeNewsNet	✓	✓	✓	✓	✓	✓	✓

how we can dynamically updating our repository in a periodical manner. The flowchart of data collection process is shown in Figure 1, which mainly consists of the collection of news contents, social context and dynamic information.

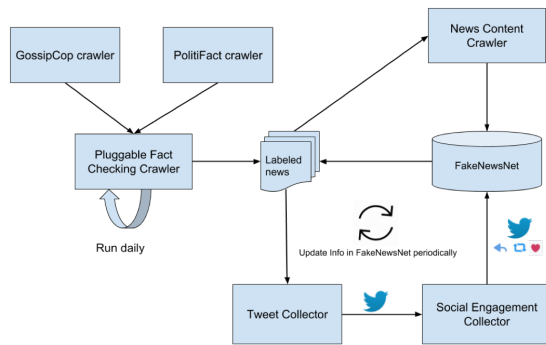


Figure 1: The flowchart of dataset construction process for FakeNewsNet. It mainly describes the collection of news content, social context and dynamic information.

News Content: To collect reliable ground truth labels for fake news, we utilize fact-checking websites to obtain news contents for fake news and true news such as *PolitiFact*¹⁰ and *GossipCop*¹¹. PolitiFact is a website operated by Tampa Bay Times, where reporters and editors from the media fact check the political news articles. PolitiFact publishes the original statement of news articles, their complete fact-check evaluation results in their website. GossipCop is a website for fact-checking entertainment stories aggregated from various media outlets. GossipCop analyzes the stories and provide truth value to each story from 0 to 10, where a rating of 0 means that the story is completely false and 10 means that the story is true. News content crawler extracts the original news source from URLs provided from the fact checking crawler. The information in news content include various information such as headline, body text, images, author information, links, etc.

Next, we introduce the details of collecting the **news content** information from PolitiFact and GossipCop. Note that the data collecting strategy we deployed can be easily

adapted to other news sources as well.

PolitiFact Crawler In PolitiFact, journalists and domain experts review the political news and provide fact-checking evaluation results to claim news articles as fake¹² and real¹³. We utilize these claims as ground truths for to fake and real news pieces.

In PolitiFact’s fact-checking evaluation result, the source URLs of the web page that published the news articles are provided, which can be used to fetch the news content related to the news articles. In some cases, the web pages of source news articles are removed and are no longer available to be collected. To tackle this problem, we make use of Google web search to identify news article that is most related to the actual news. We search the headline from the PolitiFact evaluation result on Google and the top result that does not link to fact-checking websites—in which the content from these sources are not the news content—is chosen as the news article source. One possible problem with this approach is that the quality of the news articles is based on the search results. Currently, the percentage of news pieces that can not not be collected through this strategy is about 5%. In future, selection criteria like topic distribution can be used to further enhance the quality.

GossipCop Crawler GossipCop provides rating scores on the scale of 0 to 10 to classify a news story as the degree from fake to real. From our observation, almost 90% of the stories from GossipCop have scores less than 5, which is mainly because the purpose of GossipCop is to showcase more fake stories. In order to collect true entertainment news pieces, we crawl the news articles from E! Online¹⁴, which is a well-known trusted media website for publishing entertainment news pieces. We consider all the articles from E! Online as real news sources. We collect all the news stories from GossipCop with rating scores less than 5 as the fake news stories.

Since GossipCop does not explicitly provide the URL of the source news article along with fact-checking evaluation result, so we search the news headline in Google and obtain the news source information. The headline of the GossipCop story articles are generally written in a way to reflect

¹⁰<https://www.politifact.com/>

¹¹<https://www.gossipcop.com/>

¹²available at <https://www.politifact.com/subjects/fake-news/>

¹³available at <https://www.politifact.com/truth-o-meter/rulings/true/>

¹⁴<https://www.eonline.com/>

Table 2: Statistics of the FakeNewsNet repository

Dataset Features	PolitiFact		GossipCop	
	Fake	Real	Fake	Real
Total news articles	432	624	6,048	16,817
News articles with text content	353	400	785	16,765
News articles with social engagements	342	314	4,298	2,902
News articles with both social engagements and news content	286	202	675	2,895
News articles with social engagement containing at least 1 reply	236	180	945	752
News articles with social engagement containing at least 1 like	283	219	2,911	845
News articles with social engagement containing at least 1 retweet	282	242	2,249	1,254
No. of tweets with replies	6,686	20,720	3,040	2,546
No. of tweets with likes	18,453	52,082	10,685	2,264
No. of tweets with retweets	13,226	42,059	7,614	5,025
Total no. of tweets	116,005	261,262	71,009	154,383

the actual fact and so they cannot be explicitly used to get the original news articles. So we employ different heuristic methods to formulate the search query from headline. For example, some headlines include quoted string which are exact text from the original news source. In this case, we extract the named entities through Stanfords CoreNLP tool (Manning et al. 2014) from the headline and quoted strings from the headline to form the search query. For example, in the headline *Jennifer Aniston, Brad Pitt NOT "Just Married," Despite Report*, we extract named entities including *Jennifer Aniston*, *Brad Pitt* and quoted strings including *Just Married* and form the search query as *Jennifer Aniston Brad Pitt Just Married*. In other cases, the article’s headlines are written in the negative sense to correct the false information, e.g., *"Jennifer Aniston NOT Wearing Brad Pitts Engagement Ring, Despite Report"*. So we remove negative sentiment words retrieved from SentiWordNet (Baccianella, Esuli, and Sebastiani 2010) and some hand-picked words from the headline to form the search query, e.g., *"Jennifer Aniston Wearing Brad Pitts Engagement Ring"*. We then use of the resultant search query to get the source news article. E! Online directly provides the original headline for news pieces, so we can apply the news content crawler to fetch the news content information for true news stories.

Social Context: The social engagements related to the fake and real news pieces from fact-checking websites are collected using search API provided by social media platforms such as the Twitter’s Advanced Search API¹⁵. The search queries for collecting social engagements are formed from the headlines of news articles, with special characters removed from the search query to filter out noise. After we obtain the social media posts that directly spreading news pieces, we further fetch the *second order* user behaviors towards these posts such as replies, likes, and reposts. When we obtain all the users engaging in news dissemination process, we collect all the meta data for user profiles, user posts, and the social network information.

Dynamic Information: The dynamic information indicates that we dynamically update news content and social context in a dynamic manner. We record the timestamps of

user engagements, which can be used to study how fake news pieces propagate on social media, and the topics of fake news are changing over time. Since fact-checking websites periodically update newly coming news articles, so we dynamically collect these newly added news pieces and update the FakeNewsNet repository as well. Second, we keep collecting the user engagements for all the news pieces periodically in FakeNewsNet repository such as the recent social media posts, and second order user behaviors such as replies, likes, and retweets. For example, we run the news content crawler and update Tweet collector per day. The dynamic information provides useful and comprehensive information for studying fake news problem from a temporal perspective.

4 Data Analysis

FakeNewsNet has multi-dimensional information related to news content, social context and dynamic information. In this section, we first provide some preliminary quantitative analysis to illustrate the features of FakeNewsNet. We then perform fake news detection using several state-of-the-art models to evaluate the quality of the FakeNewsNet repository. The detailed statistics of FakeNewsNet repository is illustrated in Table 2.

Assessing News Content

Since fake news attempts to spread false claims in news content, the most straightforward means of detecting it is to find clues in a news article to detect fake news. News content features describe the meta information related to a piece of news. A list of representative news content attributes include publishers, headlines, body texts, and images/videos.

First we analyze the topic distribution of fake and real news articles. From figures 2(a) and 2(b), we can observe that the fake and real news of the PolitiFact dataset are mostly related to political campaign. In case of GossipCop dataset from figures 2(c) and 2(d), we observe that the fake and real news are mostly related to gossip about relationship among celebrities. In addition, we can see the topics for fake news and real news are slightly different in general. However, for specific news, it is difficult to only use topics in the

¹⁵<https://twitter.com/search-advanced?lang=en>

content to detect fake news (Shu et al. 2017), which necessaries the need to utilize other auxiliary information such as social context.

We also explore the distribution of publishers who publish fake news on both datasets. We find out that there are in total 301 publishers publishing 432 fake news pieces, among which 191 of all publishers only publish 1 piece of fake news, and 40 publishers publish at least 2 pieces of fake news such as *washingtonpost.com* and *cnn.com*. For Gossipcop, there are total 209 publishers publishing 6,048 fake news pieces, among which 114 of all publishers only publish 1 piece of fake news, and 95 publishers publish at least 2 pieces of fake news such as *hollywoodlife.com* and *celebrityinsider.org*. The reason may be that these fact-checking websites try to identify those check-worthy breaking news events regardless of the publishers, and fake news publishers can be shut down after they were reported to publish fake news pieces.



Figure 2: The word cloud of new body text for fake and real news on PolitiFact and GossipCop.

Measuring Social Context

Social context represents the news proliferation process over time, which provides useful auxiliary information to infer the veracity of news articles. Generally, there are three major aspects of the social media context that we want to represent: user profiles, user posts, and network structures. Next, we perform an exploration study of these aspects on FakeNewsNet and introduce the potential usage of these features to help fake news detection.

User Profiles User profiles on social media have been shown to be correlated with fake news detection (Shu, Wang, and Liu 2018). Research has also shown that fake news pieces are likely to be created and spread by non-human accounts, such as social bots or cyborgs (Shu et al. 2017; Shao et al. 2017). We will illustrate some user profile features in FakeNewsNet repository.

First, we explore whether the creation time of user accounts for fake news and true news are different or not. We compute the time range of account register time with the current date and the results are shown in Figure 3. We only




Figure 3: User profile creation dates on PolitiFact




Figure 4: Comparison of bot scores on users related to fake and real news on PolitiFact dataset.

show the results for PolitiFact dataset since we have similar observations for GossipCop data. We can see that users who share real news pieces tend to have longer register time than those who share fake news. For example, around 19% of all users who share real news pieces have registered for 9 to 10 years, while only 15% for those who share fake news. The reason could be that these newly created accounts are created intentionally to spread fake news such as social bots or sybils (Shu et al. 2017).

Next, we take deeper into the user profiles and assess the social bots effects. We randomly selected 10,000 users who posted fake and real news and performed bot detection using one of the state-of-the-art bot detection algorithm Botometer (Davis et al. 2016) API¹⁶. The Botometer takes a Twitter username as an input and utilizes various features extracted from meta-data and output a probability score in $[0, 1]$, indicating how likely the user is a social bot. We set the threshold of 0.5 on the bot score returned from the Botometer results to determine bot accounts. Figure 4 shows the ratio of the bot and human users involved in tweets related to fake and real news. We can see that bots are more likely to post tweets related to fake news than real users. For example, almost 22% of users involved in fake news are bots, while only around 9% of users are predicted as bot users for real news. Similar results were observed with different thresholds on bot scores based on both datasets. This indicates that there are bots in

¹⁶<https://botometer.iuni.iu.edu/>

Table 3: The statistics of the social network of the datasets

Dataset	PolitiFact		GossipCop	
	Fake	Real	Fake	Real
# Users	214,049	700,120	99,765	69,910
# Followers	260,394,468	714,067,617	107,627,957	73,854,066
# Followees	286,205,494	746,110,345	101,790,350	75,030,435
Avg.# followers	1,216.518	1019.922	1078.815	1056.416
Avg.# followees	1,337.102	1065.689	1020.301	1073.243

Twitter for spreading fake news, which is consistent with the observation in (Shao et al. 2017).

User Posts People express their emotions or opinions towards fake news through social media posts, such as skeptical opinions, sensational reactions, etc. These features are important signals to study fake news and disinformation in general (Jin et al. 2016; Qazvinian et al. 2011).

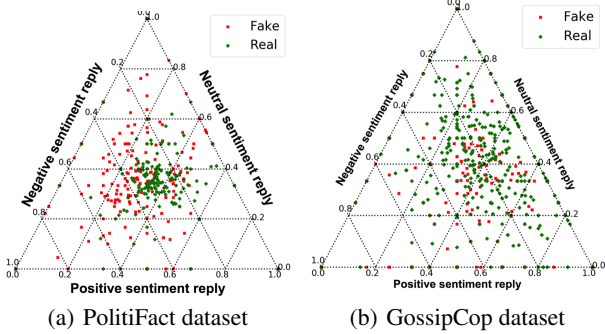


Figure 5: Ternary plot of the ratio of the positive, neutral and negative sentiment replies of the tweets related to fake and real news

We perform sentiment analysis on the replies of user posts that spreading fake news and real news using one of the state-of-the-art unsupervised prediction tool VADER (Gilbert 2014). This tool is a lexicon and rule-based sentiment analysis tool that is specifically attuned to sentiments expressed in social media¹⁷. Figure 6 shows the relationship between positive, neutral and negative replies for all news articles. For each news piece, we obtain all the replies for this news and apply VADER to predict the sentiment as positive, negative, or neutral. Then we calculate the ratio of positive, negative, and neutral replies for the news. For example, if a news piece has the sentiment distribution of replies as $[0.5, 0.5, 0.5]$, it occurs in the middle of the very center of the triangle in Figure 6. We can also see that the real news have more number of neutral replies over positive and negative replies whereas fake articles has bigger ratio of negative sentiments. In case of sentiment of the replies of the GossipCop dataset show in figure 5(b), we cannot observe any significant differences between fake and real news. This could be because of the difficulty in identifying fake and real news related to entertainment by common people.

¹⁷<https://github.com/cjhutto/vaderSentiment>

Network Structures Users tends to form different networks on social media in terms of interests, topics, and relations, which serve as the fundamental paths for information diffusion (Shu et al. 2017). Fake news dissemination processes tend to form an echo chamber cycle, highlighting the value of extracting network-based features to represent these types of network patterns for fake news detection (Del Vicario et al. 2016).




Figure 6: Distribution of the social network features like followers count and followee count of users related to fake and real news

First, we look at the social network statistics of all the users that spread fake news or real news. The social network features such as followers count and followee count can be used to estimate the scope of how the fake news can spread in social media. We plot the distribution of follower count and followee count of users in Figure 5. We can see that: i) the follower and followee count of the users generally follows power law distribution, which is commonly observed in social network structures; ii) there is a spike in the followee count distribution of both users and this is because of the restriction imposed by Twitter¹⁸ on users to have at most 5000 followees when the number of following is less than 5000.

Next, we analyze the distribution of the likes, retweets and replies of the tweets, which can help gain insights on user interaction networks related to fake and real news. Social science studies have theorized the relationship between user behaviors and their perceived beliefs on the information on social media (Kim and Dennis 2017). For example, the behaviors of likes and retweets are more emotional while replies are more rational.

We plot the ternary triangles which illustrate the ratio of replies, retweets, and likes from the second order engage-

¹⁸<https://help.twitter.com/en/using-twitter/twitter-follow-limit>

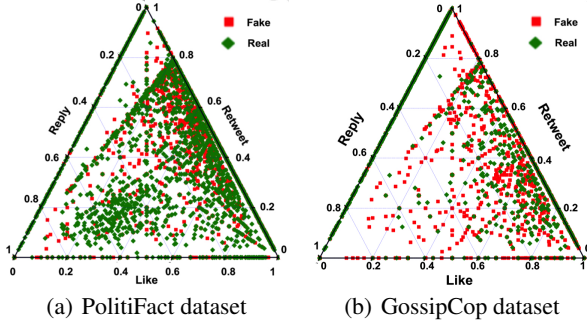


Figure 7: Ternary plot of the ratio of likes, retweet and reply of tweets related to fake and real news

ments towards the posts that spread fake news or real news pieces. From Figure 7, we observe that the: i) fake news pieces tend to have fewer replies and more retweets; ii) Real news pieces have more ratio of likes than fake news pieces, which may indicate that users are more likely to agree on real news. The differences of the distribution of user behaviors between fake news and real news could have the potential to study users’ beliefs characteristics. FakeNewsNet provides real-world datasets to understand the social factors of user engagements and underly social science as well.

Dynamic Information

Recent research has shown users’ temporal responses can be modeled using deep neural networks to help detection fake news (Ruchansky, Seo, and Liu 2017), and deep generative models can generate synthetic user engagements to help early fake news detection (Liu and Wu 2018). The dynamic information in FakeNewsNet depicts the temporal user engagements for news articles, which provides the necessary information to further study the utility of using dynamic information to detect fake news.

First, we look at whether there are differences on the temporal behaviors of users for spreading fake news and real news. We illustrate the relationship between the posting time and day of the week as in Figure 8. We can see that the time period at which the tweets related to fake news and real news posted are different. For example, in case of fake news, tweets are posted night even at odd hours between 1 am to 5 am where people are generally inactive and the density of tweets are almost the same as peak hours. This could generally because of bot accounts programmed running through out the day.

Next, we investigate if the temporal user engagements such as posts, replies, retweets, are different for fake news and real news with similar topics, in Figure 9. We can observe that: i) for fake news, there is a sudden increase in the number of retweets and it does remains constant beyond a short time whereas in case of real news, there is steady increase in the number of retweets; ii) Fake news pieces tend to receive fewer replies than real news. We also have similar observations in Table 2, and replies count for 5.76% among all Tweets for fake news, and 7.93% for real news. The differences of diffusion patterns for temporal user engagements

have the potential to determine the threshold time for early fake news detection. For example, if we can predict the sudden increase of user engagements, we should use the gathered social engagements before the time point and detect fake news accurately to limit the affect size of fake news spreading (Shu, Bernard, and Liu 2018).

Fake News Detection Performance

In this subsection, we utilize the PolitiFact and GossipCop datasets from FakeNewsNet repository to perform fake news detection task. We use 80% of data for training and 20% for testing. For evaluation metrics, we use accuracy, precision, recall and F1 score. We deployed several state-of-the-art baselines for fake news detection, as follows,


- **News Content:** News content includes the articles of the source web pages of fake and real news. We utilize the raw text features of the news articles and represent them as one-hot vectors. Then, we apply standard machine learning models including support vector machines (SVM), logistic regression (LR), Naive Bayes (NB), and CNN. For SVM, LR and NB, we used the default settings provided in the scikit-learn¹⁹ and do not tune parameters. For CNN we use the standard implementation with default setting²⁰. We also evaluate the classification on news articles using Social article fusion (SAF /S) (Shu, Mahudeswaran, and Liu) model that utilizes auto-encoder for learning features from news articles to classify a new articles as fake or real.
- **Social context:** In order to evaluate the social context, we utilize the variant of SAF model (Shu, Mahudeswaran, and Liu), i.e., SAF /A, which utilize the temporal pattern of the social engagements to detect fake news.
- **News content and social context:** To evaluate both the news and social dimension of FakeNewsNet, Social Article Fusion(SAF) model that combines SAF /S and SAF /A is used. This model uses auto-encoder with LSTM cells of 2 layers for encoder as well as decoder and also temporal pattern of the social engagements are also captures using another network of LSTM cells with 2 layers.

The experimental results are shown in Table 4. We can see that: i)For news content based methods, SAF /S perform better in terms of accuracy, recall and F1 score while logistic regression has better precision than others. SAF /A provides similar result around 66.7% accuracy as SAF /S but has higher precision.


The compared baselines models provide reasonably good performance results for the fake news detection where accuracy is mostly around 65% on PolitiFact; ii) we observe that SAF relatively achieves better accuracy than both SAF /S and SAF /A for both dataset. For example, SAF has around 5.65% and 3.68% performance improvement than SAF /S and SAF /A on PolitiFact. This indicates that social engagements can help fake news detection in addition to news articles on PolitiFact dataset.

¹⁹<http://scikit-learn.org/stable/>

²⁰<https://github.com/dennybritz/cnn-text-classification-tf>




(a) Tweets related to fake news




(b) Tweets related to real news

Figure 8: The heatmap of the day of week vs hour of tweets posted related to fake and real news



(a) Temporal user engagements of fake news



(b) Temporal user engagements of real news

Figure 9: The comparison of temporal user engagements of fake and real news: a) fake news “*TRUMP APPROVAL RATING Better than Obama and Reagan at Same Point in their Presidencies*” from June 9, 2018 to 13 June, 2018; b) real news “*President Trump in Moon Township Pennsylvania*” from March 10, 2018 to 20 March, 2018.

In summary, FakeNewsNet provides multiple dimensions of information that has the potential to benefit researchers to develop novel algorithms for fake news detection.

5 Potential Applications

FakeNewsNet contains information from multi-dimensions which could be useful for many applications. We believe FakeNewsNet would benefit the research community for studying various topics such as: (early) fake news detection, fake news evolution, fake news mitigation, malicious account detection.

Fake News Detection

One of the challenges for fake news detection is the lack of labeled benchmark dataset with reliable ground truth labels and comprehensive information space, based on which we can capture effective features and build models. FakeNewsNet can help the fake news detection task because it has reliable labels annotated by journalists and domain experts, and multi-dimension information from news content, social context and dynamic information.

First, news contents are the fundamental sources to find clues to differentiate fake news pieces. For example, study has shown that the clickbaits headlines usually can serve as a good indicator for recognizing fake news articles (Chen, Conroy, and Rubin 2015; Shu et al. 2018). In FakeNewsNet, we provide various attributes of news articles such as publishers, headlines, body texts, and images/videos.

These information can be used to extract different linguistic features (Hosseinimotlagh and Papalexakis 2018) and visual features to further build detection models for clickbaits or fake news. For example, style-based approaches try to detect fake news by capturing the manipulators in the writing style of news contents (Potthast et al. 2017; Wang 2017). In addition, Knowledge-based approaches aim to use external sources to fact-check proposed claims in news content (Shu et al. 2017). Since we directly collect news articles from fact-checking websites such as PolitiFact and GossipCop, we provide the information of detail description and explanations from the fact-checkers, which are useful for us to learn common and specific perspectives of in what aspects the fake news pieces are formed.

Second, social engagements represent the news proliferation process over time, which provides useful auxiliary information to infer the veracity of news articles. Generally, there are three major aspects of the social media context: users, generated posts, and networks. Since fake news pieces are likely to be created and spread by non-human accounts, such as social bots or cyborgs (Shao et al. 2017). Thus, capturing users profiles and characteristics by user-based features can provide useful information for fake news detection. FakeNewsNet includes all the meta data for user profiles. In addition, people express their emotions or opinions towards fake news through social media posts, such as skeptical opinions, sensational reactions, etc. We collect all the user posts for the news pieces, as well as the second engagements (see Table 2) such as reposts, comments, likes,

Table 4: Performance of the fake news detection in datasets

Model	PolitiFact				GossipCop			
	Accuracy	Precision	Recall	F1	Accuracy	Precision	Recall	F1
SVM	0.580	0.611	0.717	0.659	0.470	0.462	0.451	0.456
Logistic regression	0.642	0.757	0.543	0.633	0.822	0.897	0.722	0.799
Naive Bayes	0.617	0.674	0.630	0.651	0.704	0.735	0.765	0.798
CNN	0.629	0.807	0.456	0.583	0.703	0.789	0.623	0.699
Social Article Fusion /S	0.654	0.600	0.789	0.681	0.741	0.709	0.761	0.734
Social Article Fusion /A	0.667	0.667	0.579	0.619	0.796	0.782	0.743	0.762
Social Article Fusion	0.691	0.638	0.789	0.706	0.796	0.820	0.753	0.785

which can be utilized to extract abundant features, e.g., sentiment scores as in Figure 5, to capture fake news patterns. Moreover, fake news dissemination processes tend to form an echo chamber cycle, highlighting the value of extracting network-based features to represent these types of network patterns for fake news detection. We provide a large-scale social network of all the users involving in the news dissemination process (see Table 2).

Third, early fake news detection aims to give early alerts of fake news during the dissemination process before it reaches a broad audience (Liu and Wu 2018). Therefore early fake news detection methods are highly desirable and socially beneficial. For example, capturing the pattern of user engagements in the early phases of news diffusion could be helpful to achieve the goal of early detection. Recent approaches utilize advanced deep generative models to generate synthetic user comments to help improve fake news detection performance (Qian et al. 2018b). FakeNewsNet contains all these types of information, which provides potentials to further explore early fake news detection models.

In addition, FakeNewsNet contains two datasets of different domains, i.e., political and entertainment, which can help to study common and different patterns for fake news under different topics.

Fake News Evolution

The fake news diffusion process also has different stages in terms of peoples attentions and reactions as time goes by, resulting in a unique life cycle. For example, breaking news and in-depth news demonstrate different life cycles in social media (Castillo et al. 2014), and social media reactions can help predict future visitation patterns of news pieces accurately even at a early stage. We can have a deeper understanding of how particular stories go viral from normal public discourse by studying the fake news evolution process. First, tracking the life cycle of fake news on social media requires recording essential trajectories of fake news diffusion in general (Shao et al. 2016). Thus, FakeNewsNet has collected the related temporal user engagements which can keep track of these trajectories. Second, for specific news event, the related topics may keep changing over time and be diverse for fake news and real news. FakeNewsNet is dynamically collecting associated social engagements and allows us to perform comparison analysis (e.g., see Figure 9), and further investigate distinct temporal patterns to detect

fake news (Ruchansky, Seo, and Liu 2017). Moreover, statistical time series models such as temporal point process can be used to characterize different stages of user activities of news engagements (Farajtabar et al. 2017). FakeNewsNet enables the temporal modeling from real-world datasets, which is otherwise impossible from synthetic datasets.

Fake News Mitigation

Fake news mitigation aims to reduce the negative effects brought by fake news. During the spreading process of fake news, users play different roles such as *provenances*: the sources or originators for publishing fake news pieces; *persuaders*: who spread fake news with supporting opinions; and *clarifiers*: who propose skeptical and opposing viewpoints towards fake news and try to clarify them. Identifying key users on social media is important to mitigate the effect of fake news. For example, the provenances can help answer questions such as whether the piece of news has been modified during its propagation. In addition, its necessary to identify influential persuaders to limit the spread scope of fake news by blocking the information flow from them to their followers on social media (Shu, Bernard, and Liu 2018). FakeNewsNet provides rich information of users who were posting, liking, commenting on fake news and real news pieces (see Figure 7), which enables the exploration of identifying different types of users.

To mitigate the effect of fake news, network intervention aims to develop strategies to control the widespread dissemination of fake news before it goes viral. Two major strategies of network intervention are: i) *Influence Minimization*: minimizing the spread scope of fake news during dissemination process; ii) *Mitigation Campaign*: Limiting the impact of fake news by maximize the spread of true news. FakeNewsNet allows researchers to build diffusion network of users with temporal information, and thus can facilitate the deep understanding of minimizing the influence scopes. Furthermore, we may able to identify the fake news and real news pieces for specific event from FakeNewsNet and study the effect of mitigation campaigns in real-world datasets.

Malicious Account Detection

Studies have shown that malicious accounts that can amplify the spread of fake news include social bots, trolls, and cyborg users. Social bots are social media accounts that are controlled by a computer algorithm. Social bots can give a

false impression that information is highly popular and endorsed by many people, which enables the echo chamber effect for the propagation of fake news.

We can study the nature of the user who spread fake news and identify the characteristics of the bot account used in the fake news diffusion process through FakeNewsNet. Using the feature like the user profile meta data and the historical tweets of the user who spread fake news along with the social network one could analyze the differences in characteristics of the users to clusters the users as malicious or not. Through a preliminary study in Figure 4, we have shown that bot users are more likely to exist in fake news spreading process. Even though existing work have studied the bot detection in general, but few studies investigate the influences of social bots for fake news spreading. FakeNewsNet could potentially facilitate the study of understanding the relationship between fake news and social bots, and further explore the mutual benefits of studying fake news detection or bot detection.

6 Conclusion and Future Work

In this paper, we provide a comprehensive repository FakeNewsNet collected which contains information from news content, social context and dynamic information. We propose a principled strategy to collect relevant data from different sources. Moreover, we perform an preliminary exploration study on various features on FakeNewsNet, and demonstrate the its utility through a fake news detection tasks over several state-of-the-art baselines. FakeNewsNet has the potential to facilitate many promising research directions such as fake news detection, mitigation, evolution, malicious account detection, etc.

There are several interesting options for the future work. First, we will extend FakeNewsNet repository to other reliable news sources such as other fact checking websites or curated data collections. Second, we will improve the selection strategy used for web search results to reduce noise in the data collection process. We will also integrate FakeNewsNet repository with front-end software such as FakeNewsTracker (Shu, Mahudeswaran, and Liu), and build an end-to-end system for fake news study.

References

- [Baccianella, Esuli, and Sebastiani 2010] Baccianella, S.; Esuli, A.; and Sebastiani, F. 2010. Sentiwordnet 3.0: an enhanced lexical resource for sentiment analysis and opinion mining. In *Lrec*, volume 10, 2200–2204.
- [Castillo et al. 2014] Castillo, C.; El-Haddad, M.; Pfeffer, J.; and Stempeck, M. 2014. Characterizing the life cycle of online news stories using social media reactions. In *Proceedings of the 17th ACM conference on Computer supported cooperative work & social computing*, 211–223. ACM.
- [Chen, Conroy, and Rubin 2015] Chen, Y.; Conroy, N. J.; and Rubin, V. L. 2015. Misleading online content: Recognizing clickbait as false news. In *Proceedings of the 2015 ACM on Workshop on Multimodal Deception Detection*, 15–19. ACM.
- [Davis et al. 2016] Davis, C. A.; Varol, O.; Ferrara, E.; Flammini, A.; and Menczer, F. 2016. Botornot: A system to evaluate social bots. In *Proceedings of the 25th International Conference Companion on World Wide Web*, 273–274. International World Wide Web Conferences Steering Committee.
- [Del Vicario et al. 2016] Del Vicario, M.; Vivaldo, G.; Bessi, A.; Zollo, F.; Scala, A.; Caldarelli, G.; and Quattrociocchi, W. 2016. Echo chambers: Emotional contagion and group polarization on facebook. *Scientific reports* 6:37825.
- [Farajtabar et al. 2017] Farajtabar, M.; Yang, J.; Ye, X.; Xu, H.; Trivedi, R.; Khalil, E.; Li, S.; Song, L.; and Zha, H. 2017. Fake news mitigation via point process based intervention. *arXiv preprint arXiv:1703.07823*.
- [Gilbert 2014] Gilbert, C. H. E. 2014. Vader: A parsimonious rule-based model for sentiment analysis of social media text. In *Eighth International Conference on Weblogs and Social Media (ICWSM-14)*. Available at (20/04/16) <http://comp.social.gatech.edu/papers/icwsml4.vader.hutto.pdf>.
- [Hosseinimotlagh and Papalexakis 2018] Hosseinimotlagh, S., and Papalexakis, E. E. 2018. Unsupervised content-based identification of fake news articles with tensor decomposition ensembles.
- [Jin et al. 2016] Jin, Z.; Cao, J.; Zhang, Y.; and Luo, J. 2016. News verification by exploiting conflicting social viewpoints in microblogs. In *AAAI*, 2972–2978.
- [Kim and Dennis 2017] Kim, A., and Dennis, A. R. 2017. Says who?: How news presentation format influences perceived believability and the engagement level of social media users.
- [Liu and Wu 2018] Liu, Y., and Wu, Y.-f. B. 2018. Early detection of fake news on social media through propagation path classification with recurrent and convolutional networks. In *AAAI*.
- [Manning et al. 2014] Manning, C.; Surdeanu, M.; Bauer, J.; Finkel, J.; Bethard, S.; and McClosky, D. 2014. The stanford corenlp natural language processing toolkit. In *Proceedings of 52nd annual meeting of the association for computational linguistics: system demonstrations*, 55–60.
- [Mitra and Gilbert 2015] Mitra, T., and Gilbert, E. 2015. Credbank: A large-scale social media corpus with associated credibility annotations. In *ICWSM*, 258–267.
- [Potthast et al. 2017] Potthast, M.; Kiesel, J.; Reinartz, K.; Bevendorff, J.; and Stein, B. 2017. A stylometric inquiry into hyperpartisan and fake news. *arXiv preprint arXiv:1702.05638*.
- [Qazvinian et al. 2011] Qazvinian, V.; Rosengren, E.; Radev, D. R.; and Mei, Q. 2011. Rumor has it: Identifying misinformation in microblogs. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, 1589–1599. Association for Computational Linguistics.
- [Qian et al. 2018a] Qian, F.; Gong, C.; Sharma, K.; and Liu, Y. 2018a. Neural user response generator: Fake news detection with collective user intelligence. In *Proceedings of the Twenty-Seventh International Joint Conference on Arti-*

ficial Intelligence, IJCAI-18, 3834–3840. International Joint Conferences on Artificial Intelligence Organization.

- [Qian et al. 2018b] Qian, F.; Gong, C.; Sharma, K.; and Liu, Y. 2018b. Neural user response generator: Fake news detection with collective user intelligence. In *IJCAI*, 3834–3840.
- [Ruchansky, Seo, and Liu 2017] Ruchansky, N.; Seo, S.; and Liu, Y. 2017. Csi: A hybrid deep model for fake news detection. In *Proceedings of the 2017 ACM on Conference on Information and Knowledge Management*, 797–806. ACM.
- [Santia and Williams 2018] Santia, G. C., and Williams, J. R. 2018. Buzzface: A news veracity dataset with facebook user commentary and egos. In *ICWSM*, 531–540.
- [Shao et al. 2016] Shao, C.; Ciampaglia, G. L.; Flammini, A.; and Menczer, F. 2016. Hoaxy: A platform for tracking online misinformation. In *Proceedings of the 25th international conference companion on world wide web*, 745–750. International World Wide Web Conferences Steering Committee.
- [Shao et al. 2017] Shao, C.; Ciampaglia, G. L.; Varol, O.; Flammini, A.; and Menczer, F. 2017. The spread of fake news by social bots. *arXiv preprint arXiv:1707.07592*.
- [Shu, Bernard, and Liu 2018] Shu, K.; Bernard, H. R.; and Liu, H. 2018. Studying fake news via network analysis: Detection and mitigation. *CoRR* abs/1804.10233.
- [Shu et al. 2017] Shu, K.; Sliva, A.; Wang, S.; Tang, J.; and Liu, H. 2017. Fake news detection on social media: A data mining perspective. *ACM SIGKDD Explorations Newsletter* 19(1):22–36.
- [Shu et al. 2018] Shu, K.; Wang, S.; Le, T.; Lee, D.; and Liu, H. 2018. Deep headline generation for clickbait detection. In *ICDM*.
- [Shu, Mahudeswaran, and Liu] Shu, K.; Mahudeswaran, D.; and Liu, H. Fakenewstracker: A tool for fake news collection, detection, and visualization.
- [Shu, Wang, and Liu 2018] Shu, K.; Wang, S.; and Liu, H. 2018. Understanding user profiles on social media for fake news detection. In *2018 IEEE Conference on Multimedia Information Processing and Retrieval (MIPR)*. IEEE.
- [Tacchini et al. 2017] Tacchini, E.; Ballarin, G.; Della Vedova, M. L.; Moret, S.; and de Alfaro, L. 2017. Some like it hoax: Automated fake news detection in social networks. *arXiv preprint arXiv:1704.07506*.
- [Wang 2017] Wang, W. Y. 2017. ”liar, liar pants on fire”: A new benchmark dataset for fake news detection. *arXiv preprint arXiv:1705.00648*.